

Lending Club Interest Rate Prediction

- Description
 - This data set represents thousands of loans made through the Lending Club platform, which is a platform that allows individuals to lend to other individuals. Of course, not all loans are created equal. Someone who is a essentially a sure bet to pay back a loan will have an easier time getting a loan with a low interest rate than someone who appears to be riskier. And for people who are very risky? They may not even get a loan offer, or they may not have accepted the loan offer due to a high interest rate. It is important to keep that last part in mind, since this data set only represents loans actually made, i.e. do not mistake this data for loan applications!
- Source
 - This data comes from Lending Club (<https://www.lendingclub.com/info/statistics.action>), which provides a very large, open set of data on the people who received loans through their platform.

Walk-through of the Project

- 1. Cleansing, Preprocessing and EDA
 - Look at missing values
 - Distribution of interest rate
 - Categorical Variables -Explore categorical variables and interest rate
 - Numerical Variables -Explore numerical variables and interest rate
- 1. Feature engineering
 - Adding more variables
 - Scaling & Getting dummy
 - Feature selection(Lasso CV)
- 1. Model
 - Random Forest
 - XGBoost

Import data

In [155]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer
from string import ascii_letters
from sklearn.model_selection import RepeatedKFold
from sklearn.model_selection import cross_val_score
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.preprocessing import MinMaxScaler, StandardScaler
from sklearn.preprocessing import Normalizer
from sklearn.model_selection import train_test_split
%matplotlib inline
```

In [2]:

```
data=pd.read_csv("loans_full_schema.csv")
```

In [3]:

```
data.head()
```

Out[3]:

	emp_title	emp_length	state	homeownership	annual_income	verified_income	debt_to_in
0	global config engineer	3.0	NJ	MORTGAGE	90000.0	Verified	
1	warehouse office clerk	10.0	HI	RENT	40000.0	Not Verified	
2	assembly	3.0	WI	RENT	40000.0	Source Verified	
3	customer service	1.0	PA	RENT	30000.0	Not Verified	
4	security supervisor	10.0	CA	RENT	35000.0	Verified	

5 rows × 55 columns

Cleansing, Preprocessing and EDA

In [4]:

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 55 columns):
#   Column                                     Non-Null Count  Dtype
---  -
#   Column                                     Non-Null Count  Dtype
```

0	emp_title	9167	non-null	object
1	emp_length	9183	non-null	float64
2	state	10000	non-null	object
3	homeownership	10000	non-null	object
4	annual_income	10000	non-null	float64
5	verified_income	10000	non-null	object
6	debt_to_income	9976	non-null	float64
7	annual_income_joint	1495	non-null	float64
8	verification_income_joint	1455	non-null	object
9	debt_to_income_joint	1495	non-null	float64
10	delinq_2y	10000	non-null	int64
11	months_since_last_delinq	4342	non-null	float64
12	earliest_credit_line	10000	non-null	int64
13	inquiries_last_12m	10000	non-null	int64
14	total_credit_lines	10000	non-null	int64
15	open_credit_lines	10000	non-null	int64
16	total_credit_limit	10000	non-null	int64
17	total_credit_utilized	10000	non-null	int64
18	num_collections_last_12m	10000	non-null	int64
19	num_historical_failed_to_pay	10000	non-null	int64
20	months_since_90d_late	2285	non-null	float64
21	current_accounts_delinq	10000	non-null	int64
22	total_collection_amount_ever	10000	non-null	int64
23	current_installment_accounts	10000	non-null	int64
24	accounts_opened_24m	10000	non-null	int64
25	months_since_last_credit_inquiry	8729	non-null	float64
26	num_satisfactory_accounts	10000	non-null	int64
27	num_accounts_120d_past_due	9682	non-null	float64
28	num_accounts_30d_past_due	10000	non-null	int64
29	num_active_debit_accounts	10000	non-null	int64
30	total_debit_limit	10000	non-null	int64
31	num_total_cc_accounts	10000	non-null	int64
32	num_open_cc_accounts	10000	non-null	int64
33	num_cc_carrying_balance	10000	non-null	int64
34	num_mort_accounts	10000	non-null	int64
35	account_never_delinq_percent	10000	non-null	float64
36	tax_liens	10000	non-null	int64
37	public_record_bankrupt	10000	non-null	int64
38	loan_purpose	10000	non-null	object
39	application_type	10000	non-null	object
40	loan_amount	10000	non-null	int64
41	term	10000	non-null	int64
42	interest_rate	10000	non-null	float64
43	installment	10000	non-null	float64
44	grade	10000	non-null	object
45	sub_grade	10000	non-null	object
46	issue_month	10000	non-null	object
47	loan_status	10000	non-null	object
48	initial_listing_status	10000	non-null	object
49	disbursement_method	10000	non-null	object
50	balance	10000	non-null	float64
51	paid_total	10000	non-null	float64
52	paid_principal	10000	non-null	float64
53	paid_interest	10000	non-null	float64
54	paid_late_fees	10000	non-null	float64

dtypes: float64(17), int64(25), object(13)

memory usage: 4.2+ MB

10000 sample size with 55columns.

```
In [5]: df_miss=data.isnull().sum()/len(data)*100
df_miss=pd.DataFrame(df_miss,columns=['percentage'])
```

Difficult for to fill in the null values because emp_title are categorical vairables.Too much category.Just delete this colum

```
In [6]: ##delete missing data>50%
df_miss[df_miss['percentage']>50]
```

```
Out[6]:
```

	percentage
annual_income_joint	85.05
verification_income_joint	85.45
debt_to_income_joint	85.05
months_since_last_delinq	56.58
months_since_90d_late	77.15

```
In [7]: ##drop vairables missing values percentage>50
data_new=data.drop(df_miss[df_miss['percentage']>50].index,axis=1)
```

```
In [8]: data_new.columns
```

```
Out[8]: Index(['emp_title', 'emp_length', 'state', 'homeownership', 'annual_income',
               'verified_income', 'debt_to_income', 'delinq_2y',
               'earliest_credit_line', 'inquiries_last_12m', 'total_credit_lines',
               'open_credit_lines', 'total_credit_limit', 'total_credit_utilized',
               'num_collections_last_12m', 'num_historical_failed_to_pay',
               'current_accounts_delinq', 'total_collection_amount_ever',
               'current_installment_accounts', 'accounts_opened_24m',
               'months_since_last_credit_inquiry', 'num_satisfactory_accounts',
               'num_accounts_120d_past_due', 'num_accounts_30d_past_due',
               'num_active_debit_accounts', 'total_debit_limit',
               'num_total_cc_accounts', 'num_open_cc_accounts',
               'num_cc_carrying_balance', 'num_mort_accounts',
               'account_never_delinq_percent', 'tax_liens', 'public_record_bankrupt',
               'loan_purpose', 'application_type', 'loan_amount', 'term',
               'interest_rate', 'installment', 'grade', 'sub_grade', 'issue_month',
               'loan_status', 'initial_listing_status', 'disbursement_method',
               'balance', 'paid_total', 'paid_principal', 'paid_interest',
               'paid_late_fees'],
              dtype='object')
```

Dealing with categorical variables

```
In [9]: cat_cols = data_new.select_dtypes(include=("object"))
```

```
In [10]: cat_cols.columns
```

```
Out[10]: Index(['emp_title', 'state', 'homeownership', 'verified_income',
               'loan_purpose', 'application_type', 'grade', 'sub_grade', 'issue_month',
               'loan_status', 'initial_listing_status', 'disbursement_method'],
              dtype='object')
```

```
In [11]: cat_cols['emp_title'].nunique()
```

```
Out[11]: 4741
```

```
In [12]: cat_cols['state'].nunique()
```

```
Out[12]: 50
```

```
In [13]: cat_cols['emp_title'].value_counts()
```

```
Out[13]: manager                218
owner                204
teacher              201
driver              123
sales                97
...
cooler service technician    1
its5                          1
nursing supervisor          1
partner physician           1
brand ambassador/promotional model  1
Name: emp_title, Length: 4741, dtype: int64
```

```
In [14]: ## Too many unique values for employment title, which is low
cat_cols=cat_cols.drop(['emp_title'],axis=1)
```

```
In [15]: for i in cat_cols.columns.tolist():
          print(data_new[i].value_counts())
```

```
CA    1330
TX     806
NY     793
FL     732
```

IL	382
NJ	338
OH	338
GA	334
NC	299
PA	298
VA	261
AZ	255
MD	247
MI	245
MA	237
WA	235
CO	235
CT	181
IN	178
TN	167
MO	159
MN	159
NV	158
SC	145
OR	130
WI	128
AL	122
KY	97
LA	96
KS	89
OK	81
MS	72
AR	70
WV	68
UT	61
NE	56
RI	53
NH	47
NM	43
ID	38
HI	35
AK	33
ME	26
MT	24
DE	24
VT	23
SD	20
WY	19
DC	19
ND	14

Name: state, dtype: int64

MORTGAGE	4789
RENT	3858
OWN	1353

Name: homeownership, dtype: int64

Source Verified	4116
Not Verified	3594
Verified	2290

Name: verified_income, dtype: int64

debt_consolidation	5144
credit_card	2249

```

other          914
home_improvement 680
major_purchase 303
medical        162
house          151
car            131
small_business 125
moving         69
vacation       62
renewable_energy 10
Name: loan_purpose, dtype: int64
individual     8505
joint         1495
Name: application_type, dtype: int64
B             3037
C             2653
A             2459
D             1446
E              335
F              58
G              12
Name: grade, dtype: int64
B1            647
B2            638
B5            631
C1            597
B4            581
C2            572
A4            568
B3            540
A3            504
C4            503
C3            501
A5            485
A2            480
C5            480
A1            422
D2            323
D1            319
D3            311
D5            262
D4            231
E5            91
E3            73
E4            73
E2            57
E1            41
F1            31
G1            11
F2             9
F3             9
F4             5
F5             4
G4             1
Name: sub_grade, dtype: int64
Mar-2018      3617
Jan-2018      3395

```

```

Feb-2018      2988
Name: issue_month, dtype: int64
Current      9375
Fully Paid    447
In Grace Period    67
Late (31-120 days)    66
Late (16-30 days)    38
Charged Off      7
Name: loan_status, dtype: int64
whole      8206
fractional   1794
Name: initial_listing_status, dtype: int64
Cash      9284
DirectPay   716
Name: disbursement_method, dtype: int64

```

```
In [16]: from scipy.stats import spearmanr
```

```
In [17]: ## Not much correlaiton between state and interest rate
data_new['state'].corr(data_new['interest_rate'],method='spearman')
```

```
Out[17]: 0.0011210910173769094
```

```
In [18]: cat_cols=cat_cols.drop(['state'],axis=1)
```

the correlation is not strong and delete the state columns

```
In [19]: for i in cat_cols.columns.tolist():
          cor=data_new[i].corr(data_new['interest_rate'],method='spearman')
          print(i,cor)
```

```

homeownership 0.08748952903152749
verified_income 0.24571897146283708
loan_purpose 0.0476652973560531
application_type 0.053857860348730205
grade 0.9666157034979389
sub_grade 0.9981746241279144
issue_month -0.03445875962956296
loan_status 0.08271675152667414
initial_listing_status -0.10904837269519452
disbursement_method -0.18045601139148257

```

```
In [222... for i in cat_cols.columns.tolist():
            cor=data_new[i].corr(data_new['grade'],method='spearman')
            print(i,cor)
```



```
homeownership 0.07812752527000763
verified_income 0.23479181403496524
loan_purpose 0.05090817031653836
application_type 0.05495150790610031
grade 0.9999999999999999
sub_grade 0.968431774166524
issue_month -0.012100618280686318
loan_status 0.07546741175826262
initial_listing_status -0.11377718580877846
disbursement_method -0.17106527952418277
```

```
In [21]: data_new['interest_rate'].round().value_counts()
```

```
Out[21]: 10.0    1178
         14.0    1073
         7.0     1072
         9.0     647
        12.0     631
        13.0     597
        11.0     581
        15.0     503
         8.0     485
         6.0     482
        16.0     480
         5.0     422
        18.0     322
        17.0     319
        19.0     311
        20.0     230
        21.0     172
        26.0     126
        22.0      90
        25.0      73
        24.0      57
        23.0      41
        27.0      38
        29.0      31
        31.0      21
        30.0      18
Name: interest_rate, dtype: int64
```

```
In [224... data_new.describe()
```

Out [224]...

	emp_length	annual_income	debt_to_income	delinq_2y	earliest_credit_line	inquiries
count	9183.000000	1.000000e+04	9976.000000	10000.000000	10000.000000	1
mean	5.930306	7.922215e+04	19.308192	0.21600	2001.29000	
std	3.703734	6.473429e+04	15.004851	0.68366	7.79551	
min	0.000000	0.000000e+00	0.000000	0.00000	1963.00000	
25%	2.000000	4.500000e+04	11.057500	0.00000	1997.00000	
50%	6.000000	6.500000e+04	17.570000	0.00000	2003.00000	
75%	10.000000	9.500000e+04	25.002500	0.00000	2006.00000	
max	10.000000	2.300000e+06	469.090000	13.00000	2015.00000	

8 rows × 42 columns

Many variables containing outliers and missing values

Distribution of interest rate

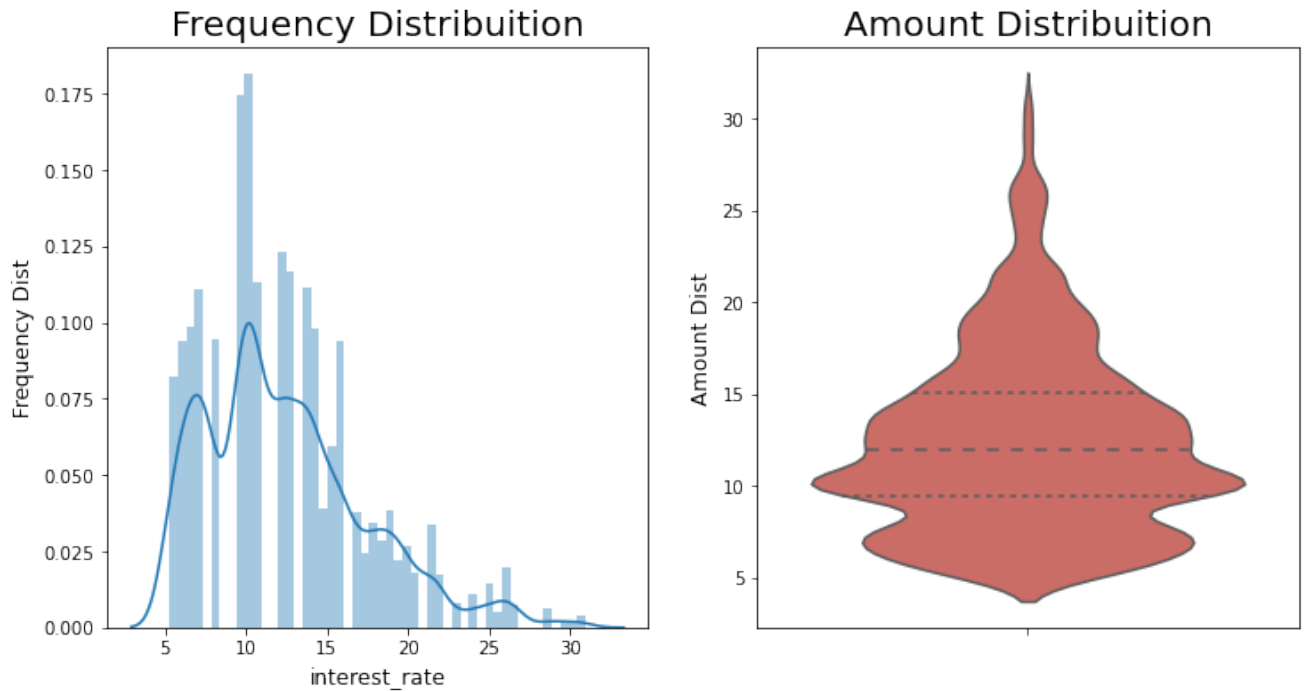
In [22]:

```
plt.figure(figsize=(12,6))
plt.subplot(121)
g = sns.distplot(data_new["interest_rate"])
g.set_xlabel("interest_rate", fontsize=12)
g.set_ylabel("Frequency Dist", fontsize=12)
g.set_title("Frequency Distribution", fontsize=20)
plt.subplot(122)
g1 = sns.violinplot(y="interest_rate", data=data_new,
                    inner="quartile", palette="hls")
g1.set_xlabel("", fontsize=12)
g1.set_ylabel("Amount Dist", fontsize=12)
g1.set_title("Amount Distribution", fontsize=20)

plt.show()
```

/Users/yanzhenlei/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)




In [24]:

```
## The distribution of interest rate is right-skewed, we should log transform
#Exploring the Int_rate
data_new['log_int_rate']=np.log(data_new["interest_rate"])
plt.figure(figsize=(12,6))
plt.subplot(211)
g = sns.distplot(data_new['log_int_rate'])
g.set_xlabel("", fontsize=12)
g.set_ylabel("Distribution", fontsize=12)
g.set_title("Int Rate Log distribution", fontsize=20)
plt.show()
```

```
/Users/yanzhenlei/opt/anaconda3/lib/python3.8/site-packages/seaborn/distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
```

```
warnings.warn(msg, FutureWarning)
```



If we are transform the distribution of interest rate, we can apply linear regression


Explore categorical variables and interest rate

Grades and Subgrades

In [96]:

```
##Grade& interest rate
plt.figure(figsize=(12,12))
plt.subplot(211)
sns.boxplot(x="grade", y="interest_rate", data=data_new,palette="hls", hue="a")
plt.subplot(212)
sns.boxenplot(x="sub_grade", y="interest_rate", data=data_new,palette="hls")
```


Out[96]: <AxesSubplot:xlabel='sub_grade', ylabel='interest_rate'>



It is obvious that grade had high correlation with interest rate. Higher grade, higher interest rate.

```
In [25]: ##Loan_status& interest rate
sns.displot(data=data_new, x="interest_rate", hue='loan_status', kind='kde')
```

Out[25]: <seaborn.axisgrid.FacetGrid at 0x7ffc8ae32eb0>



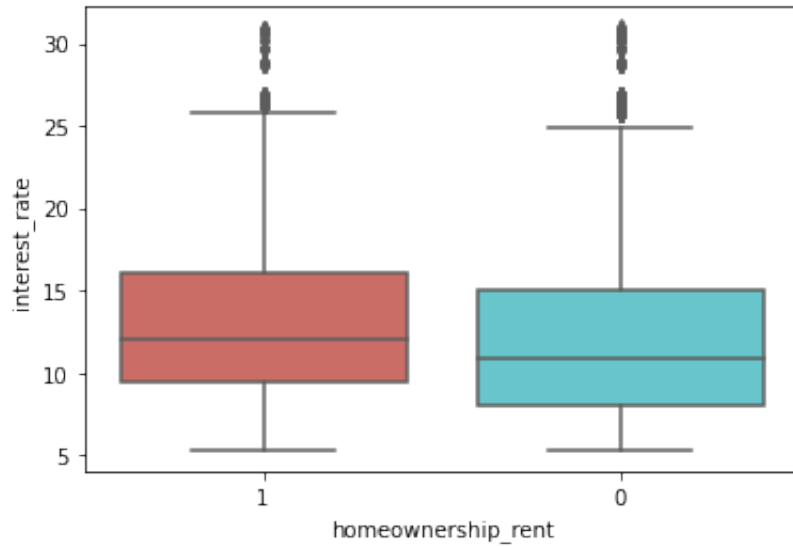
```
In [28]: ##Homeownership& interest rate
data_new['homeownership'].value_counts()
```

```
Out[28]: MORTGAGE    4789
RENT          3858
OWN           1353
Name: homeownership, dtype: int64
```

```
In [30]: data_new['homeownership_rent']=np.where(data_new['homeownership']=='RENT',1,0)
```

```
In [33]: fig= plt.figure(figsize=(24,10))
sns.boxplot(x="homeownership_rent", y="interest_rate", data=data_new,palette=
```

Out[33]: <AxesSubplot:xlabel='homeownership_rent', ylabel='interest_rate'>



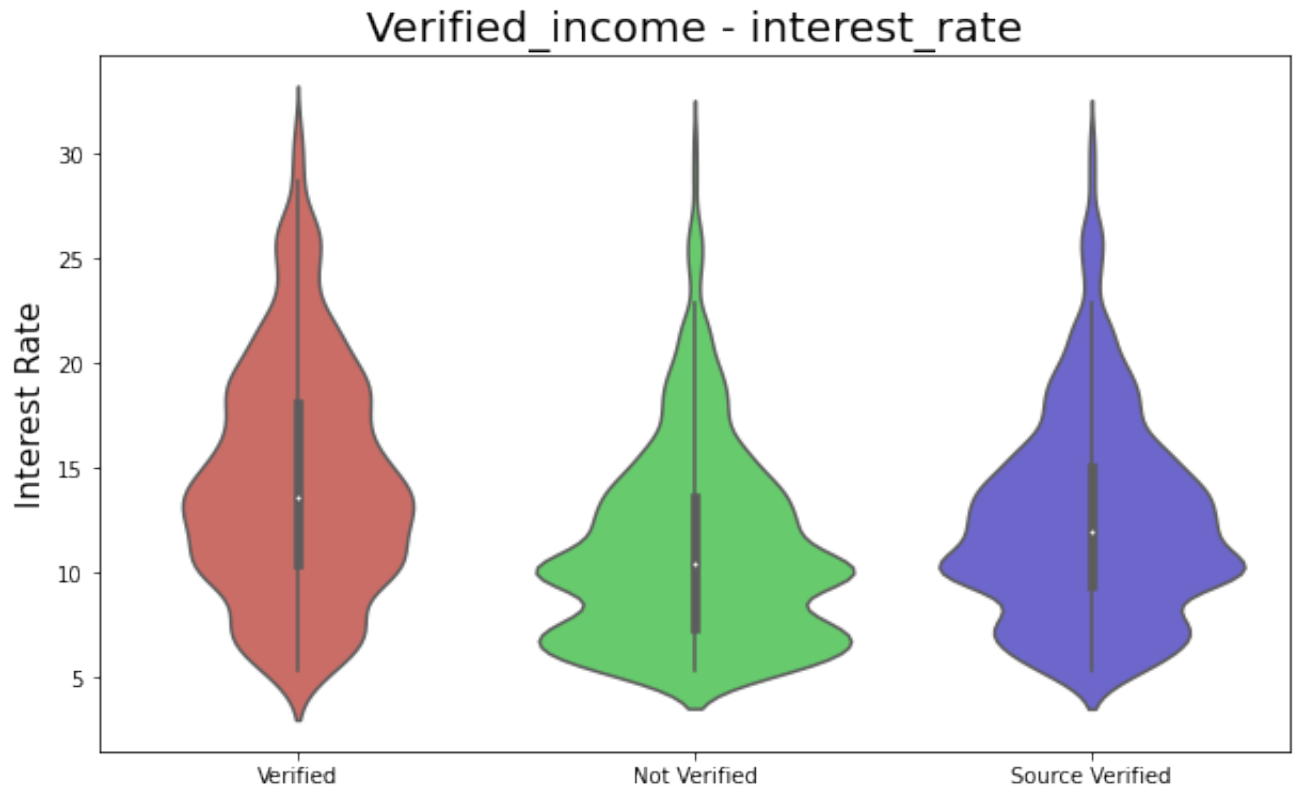
```
In [34]: ##Homeownership& interest rate
data_new['verified_income'].value_counts()
```

```
Out[34]: Source Verified      4116
Not Verified    3594
Verified        2290
Name: verified_income, dtype: int64
```

```
In [35]: plt.figure(figsize = (10,6))

g = sns.violinplot(x="verified_income",y="interest_rate",data=data_new,
                  kind="violin",
                  split=True,palette="hls")
g.set_title("Verified_income - interest_rate", fontsize=20)
g.set_xlabel("", fontsize=15)
g.set_ylabel("Interest Rate", fontsize=15)

plt.show()
```



higher correlation with grade, maintain the columns

All categorical variables are bearable number of categories

```
In [72]: df_cat = pd.get_dummies(cat_cols, columns=cat_cols.columns.tolist(), drop_first=True)
df_cat
```

Numerical columns processing

```
In [39]: # Subset numeric features: numeric_cols
numeric_cols = data_new.select_dtypes(include=[np.number])
# Iteratively impute
imp_iter = IterativeImputer(max_iter=5, sample_posterior=True, random_state=1)
loans_imp_iter = imp_iter.fit_transform(numeric_cols)
# Convert returned array to DataFrame
loans_imp_iterDF = pd.DataFrame(loans_imp_iter, columns=numeric_cols.columns)
```


```
In [40]: numeric_cols.columns
```



```
Out[40]: Index(['emp_length', 'annual_income', 'debt_to_income', 'delinq_2y',  
              'earliest_credit_line', 'inquiries_last_12m', 'total_credit_lines',  
              'open_credit_lines', 'total_credit_limit', 'total_credit_utilized',  
              'num_collections_last_12m', 'num_historical_failed_to_pay',  
              'current_accounts_delinq', 'total_collection_amount_ever',  
              'current_installment_accounts', 'accounts_opened_24m',  
              'months_since_last_credit_inquiry', 'num_satisfactory_accounts',  
              'num_accounts_120d_past_due', 'num_accounts_30d_past_due',  
              'num_active_debit_accounts', 'total_debit_limit',  
              'num_total_cc_accounts', 'num_open_cc_accounts',  
              'num_cc_carrying_balance', 'num_mort_accounts',  
              'account_never_delinq_percent', 'tax_liens', 'public_record_bankrupt',  
              'loan_amount', 'term', 'interest_rate', 'installment', 'balance',  
              'paid_total', 'paid_principal', 'paid_interest', 'paid_late_fees',  
              'log_int_rate', 'homeownership_rent'],  
             dtype='object')
```

```
In [41]: ##correlations  
sns.set(style="white")  
  
corr = loans_imp_iterDF.corr()  
  
mask = np.zeros_like(corr, dtype=np.bool)  
mask[np.triu_indices_from(mask)] = True  
  
f, ax = plt.subplots(figsize=(11, 9))  
  
cmap = sns.diverging_palette(220, 10, as_cmap=True)  
  
sns.heatmap(corr, mask=mask, cmap=cmap, vmax=.3, center=0,  
           square=True, linewidths=.5, cbar_kws={"shrink": .5})
```

Out[41]: <AxesSubplot:>




```
In [81]: X_num=loans_imp_iterDF.drop(['interest_rate'],axis=1)
for i in X_num.columns.tolist():
    cor=loans_imp_iterDF[i].corr(loans_imp_iterDF['interest_rate'])
    print(i,cor)
```

```
emp_length -0.022864260617945827
annual_income -0.09958433559316995
debt_to_income 0.14013187435197783
delinq_2y 0.09045622123463228
earliest_credit_line 0.10363461600403122
inquiries_last_12m 0.13119296831174532
total_credit_lines -0.04443485546943214
open_credit_lines -0.012773330697410155
total_credit_limit -0.1304965801657026
total_credit_utilized 0.03152743152587917
num_collections_last_12m 0.02990136827787591
num_historical_failed_to_pay 0.03756269731790593
current_accounts_delinq 0.007184083055156142
total_collection_amount_ever 0.013532163648097315
current_installment_accounts 0.047131258024430724
accounts_opened_24m 0.120656747256795
months_since_last_credit_inquiry -0.08684992876486065
num_satisfactory_accounts -0.013329575570643909
num_accounts_120d_past_due -0.012513901637688854
num_accounts_30d_past_due 0.007184083055156142
num_active_debit_accounts 0.0268002604908547
total_debit_limit -0.254225259596709
num_total_cc_accounts -0.07123798773636468
num_open_cc_accounts -0.028891174410642345
num_cc_carrying_balance 0.0817406949144792
num_mort_accounts -0.13835135211087057
account_never_delinq_percent -0.12470193063277377
tax_liens 0.016547306573736924
public_record_bankrupt 0.047476574571926194
loan_amount 0.06452688502519918
installment 0.09881137993313831
balance 0.067569664942812
paid_total 0.06533871565081698
paid_principal -0.0019527818264753116
paid_interest 0.514507529964494
paid_late_fees 0.02310998704692934
log_int_rate 0.9720199909247791
homeownership_rent 0.07883478471311552
```

In [43]:

```
plt.plot(data_new['term'],data_new['interest_rate'],'bo')
plt.show()
```

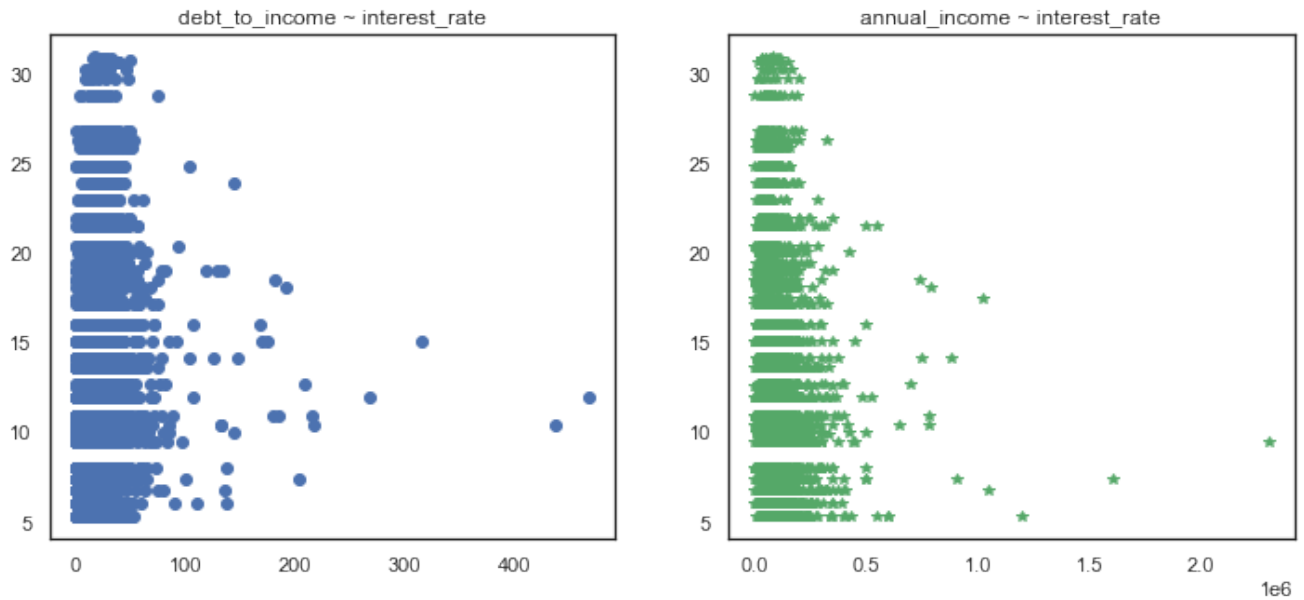


```
In [45]: data_new['term_36']=data_new['term'][data_new['term']==36]
data_new['term_60']=data_new['term'][data_new['term']==60]
```

```
In [46]: data_new['term_36']=np.where(data_new['term_36']==36,1,0)
data_new['term_60']=np.where(data_new['term_60']==60,1,0)
```

```
In [47]: loans_imp_iterDF=loans_imp_iterDF.drop(['term'],axis=1)
```

```
In [221... plt.figure(figsize=(12,5))
plt.subplot(121)
plt.plot(data_new['debt_to_income'],data_new['interest_rate'],'bo')
plt.title('debt_to_income ~ interest_rate')
plt.subplot(122)
plt.plot(data_new['annual_income'],data_new['interest_rate'],'g*')
plt.title('annual_income ~ interest_rate')
plt.show()
```



```
In [82]: df_for_RF=pd.concat([df_cat,loans_imp_iterDF],axis=1)
```

```
In [83]: len(df_for_RF.columns)
```

```
Out[83]: 101
```

```
In [84]: df_total['log_int_rate'].isnull().sum()
```

```
Out[84]: 0
```

```
df_total.to_csv('df_total.csv')
```

DATA Scaling

```
In [127... #normalise the data
X_train, y_train = Train.drop(['log_int_rate','interest_rate'], axis=1), Train['interest_rate']
X_test, y_test = Test.drop(['log_int_rate','interest_rate'], axis=1), Test['interest_rate']
scaler = StandardScaler().fit(X_train)
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
df_for_RF=pd.merge(X_train,
```

Feature selection

Too much variables for 10000 samples

In []:

df_for_RF

In [90]:

from sklearn.ensemble import RandomForestClassifier

In [101]:

df_for_RF.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Columns: 101 entries, homeownership_OWEN to homeownership_rent
dtypes: float64(39), uint8(62)
memory usage: 3.6 MB

In [109]:

size = df_for_RF.shape[0]

Train, Test_new = train_test_split(df_for_RF, test_size= 0.3, random_state= 1)

CV, Test_new = train_test_split(Test_new, test_size=0.5, random_state = 1)

print(Train.shape, CV.shape, Test_new.shape)

(7000, 101) (1500, 101) (1500, 101)

In [115]:

CV

Out[115]:

	homeownership_OWEN	homeownership_RENTE	verified_income_Source Verified	verified_income_V
4940	0	0	1	
7276	0	0	0	
7829	0	0	1	
9466	1	0	1	
8477	0	0	0	
...	
401	0	0	1	
6899	0	0	0	
5938	0	1	1	
3632	0	0	0	
2247	0	0	0	

1500 rows x 101 columns

```
In [133... X_train, y_train = Train.drop(['log_int_rate', 'interest_rate'], axis=1), Train
X_test, y_test = Test_new.drop(['log_int_rate', 'interest_rate'], axis=1), Test_new
CV_x, CV_y = CV.drop(['log_int_rate', 'interest_rate'], axis=1), CV.interest_rate
x_col=X_train.columns
```

```
In [134... X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
CV_x=scaler.transform(CV_x)
```

```
In [135... Test.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 1500 entries, 5309 to 5509
Columns: 101 entries, homeownership_OWN to homeownership_rent
dtypes: float64(39), uint8(62)
memory usage: 559.6 KB
```

```
In [136... CV_x
```

```
Out[136... array([[ -0.39384707, -0.79743102,  1.19570561, ..., -0.50967092,
        -0.06058442, -0.79743102],
       [ -0.39384707, -0.79743102, -0.83632626, ...,  0.22273732,
        -0.06058442, -0.79743102],
       [ -0.39384707, -0.79743102,  1.19570561, ...,  1.4397422 ,
        -0.06058442, -0.79743102],
       ...,
       [ -0.39384707,  1.25402696,  1.19570561, ...,  0.26850925,
        -0.06058442,  1.25402696],
       [ -0.39384707, -0.79743102, -0.83632626, ..., -0.53002038,
        -0.06058442, -0.79743102],
       [ -0.39384707, -0.79743102, -0.83632626, ..., -0.75516619,
        -0.06058442, -0.79743102]])
```

```
In [137... from sklearn.linear_model import LassoCV
```

In [140...

```

modellasso = LassoCV(alphas = [1, 0.1, 0.001, 0.0001, 10, 1000]).fit(X_train,
lassopred = modellasso.predict(CV_x)
print("RMSE of Lasso: ", np.sqrt(mean_squared_error(lassopred, CV_y)))

coeff = modellasso.coef_

x = list(x_col)
x_pos = [i for i, _ in enumerate(x)]

plt.figure(figsize = (10,40))
plt.barh(x_pos, coeff, color='green')
plt.ylabel("Features -->")
plt.xlabel("Coefficients -->")
plt.title("Coefficients from Lasso")
plt.yticks(x_pos, x)

plt.show()

```

/Users/yanzhenlei/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_coordinate_descent.py:526: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 47.482425143210094, tolerance: 13.976402134792863

model = cd_fast.enet_coordinate_descent_gram(
/Users/yanzhenlei/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_coordinate_descent.py:526: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 47.74899457664486, tolerance: 14.105763850998242

model = cd_fast.enet_coordinate_descent_gram(
/Users/yanzhenlei/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_coordinate_descent.py:526: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 48.38152692615847, tolerance: 14.244188974555362

model = cd_fast.enet_coordinate_descent_gram(
/Users/yanzhenlei/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_coordinate_descent.py:526: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 40.249599592524504, tolerance: 13.933095255312502

model = cd_fast.enet_coordinate_descent_gram(
/Users/yanzhenlei/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_coordinate_descent.py:526: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 48.32462362730462, tolerance: 14.07528496898393


model = cd_fast.enet_coordinate_descent_gram(
/Users/yanzhenlei/opt/anaconda3/lib/python3.8/site-packages/sklearn/linear_model/_coordinate_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 117.7406962564507, tolerance: 17.58403962490858

model = cd_fast.enet_coordinate_descent(
RMSE of Lasso: 0.3729692584224117

Coefficients from Lasso



Features -->	homeownership_rent	
	paid_late_fees	
	paid_interest	
	paid_principal	
	paid_total	
	balance	
	installment	
	loan_amount	
	public_record_bankrupt	
	tax_liens	
	account_never_delinq_percent	
	num_mort_accounts	
	num_cc_carrying_balance	
	num_open_cc_accounts	
	num_total_cc_accounts	
	total_debit_limit	
	num_active_debit_accounts	
	num_accounts_30d_past_due	
	num_accounts_120d_past_due	
	num_satisfactory_accounts	
	months_since_last_credit_inquiry	
	accounts_opened_24m	
	current_installment_accounts	
	total_collection_amount_ever	
	current_accounts_delinq	
	num_historical_failed_to_pay	
	num_collections_last_12m	
	total_credit_utilized	
	total_credit_limit	
	open_credit_lines	
	total_credit_lines	
	inquiries_last_12m	
	earliest_credit_line	
	delinq_2y	
	debt_to_income	
	annual_income	
	emp_length	
	disbursement_method_DirectPay	
	initial_listing_status_whole	
	loan_status_Late (31-120 days)	
	loan_status_Late (16-30 days)	
	loan_status_In Grace Period	
	loan_status_Fully Paid	
	loan_status_Current	
	issue_month_Mar-2018	
	issue_month_Jan-2018	
	sub_grade_G4	
	sub_grade_G1	
	sub_grade_F5	
	sub_grade_F4	
	sub_grade_F3	



Coefficients -->

```
In [142... data_col=pd.DataFrame(columns=['name','coeff'])
co=list(coeff)
```

```
In [143... data_col['name']=x_col
data_col['coeff']=co
```

```
In [144... data_col.sort_values(by='coeff',ascending=False)
```

```
Out[144...
```

	name	coeff
18	grade_D	4.557579
19	grade_E	3.333528
17	grade_C	3.205376
20	grade_F	1.692306
30	sub_grade_B5	1.067477
...
53	issue_month_Jan-2018	-0.004499
68	total_credit_lines	-0.005607
31	sub_grade_C1	-0.005758
41	sub_grade_E1	-0.062488
36	sub_grade_D1	-0.175357

99 rows x 2 columns

```
In [145... data_col.count()
```

```
Out[145... name      99
coeff      99
dtype: int64
```

```
In [183... col=data_col[abs(data_col['coeff'])>0.5]
```

```
In [184... col.count()
```

```
Out[184... name      14  
          coeff     14  
          dtype: int64
```

```
In [185... col['name'].tolist()
```

```
Out[185... ['grade_B',  
            'grade_C',  
            'grade_D',  
            'grade_E',  
            'grade_F',  
            'grade_G',  
            'sub_grade_A5',  
            'sub_grade_B2',  
            'sub_grade_B3',  
            'sub_grade_B4',  
            'sub_grade_B5',  
            'sub_grade_C4',  
            'sub_grade_C5',  
            'sub_grade_D5']
```

Random Forest

- Advantages of random forest
 - It can perform both regression and classification tasks.
 - A random forest produces good predictions that can be understood easily.
 - It can handle large datasets efficiently.
 - The random forest algorithm provides a higher level of accuracy in predicting outcomes over the decision tree algorithm.
- Disadvantages of random forest
 - When using a random forest, more resources are required for computation.
 - It consumes more time compared to a decision tree algorithm.

```
In [186... features= df_for_RF[col['name'].tolist()]  
labels=df_total['interest_rate']
```

```
In [187... train_features, test_features, train_labels, test_labels = train_test_split(f
```

```
In [188... print('Training Features Shape:', train_features.shape)  
print('Testing Features Shape:', test_features.shape)  
print('Training Labels Shape:', train_labels.shape)  
print('Testing Labels Shape:', test_labels.shape)
```

```
Training Features Shape: (7000, 14)
Testing Features Shape: (3000, 14)
Training Labels Shape: (7000,)
Testing Labels Shape: (3000,)
```

```
In [189... from sklearn.ensemble import RandomForestRegressor
```

```
In [209... # Instantiate model (Using Default 10 Estimators)
rf = RandomForestRegressor(n_estimators= 10, random_state=42)

# Using Evaluation Function on our First Model

rf.fit(train_features, train_labels)

y_test_pred = rf.predict(test_features)
```

```
In [210... y_test_pred.shape
```

```
Out[210... (3000,)
```

```
In [211... # Mean squared error
from sklearn import metrics
print("Mean squared error: %.2f" % mean_squared_error(test_labels, y_test_pre
print('Mean Absolute Percentage Error (MAPE):', round(metrics.mean_absolute_p
print('Accuracy:', round(100*(1 - metrics.mean_absolute_percentage_error(test
```

```
Mean squared error: 0.36
Mean Absolute Percentage Error (MAPE): 3.55
Accuracy: 96.45
```

Randomforest tuning

For model performance improvement, We should use parameter tuning

- `n_estimators`
 - The `n_estimators` parameter specifies the number of trees in the forest of the model. The default value for this parameter is 10, which means that 10 different decision trees will be constructed in the random forest.
- `max_depth`
 - The `max_depth` parameter specifies the maximum depth of each tree. The default value for `max_depth` is `None`, which means that each tree will expand until every leaf is pure. A pure leaf is one where all of the data on the leaf comes from the same class.
- `min_samples_split`
 - The `min_samples_split` parameter specifies the minimum number of samples required to split an internal leaf node. The default value for this parameter is 2, which means that an internal node must have at least two samples before it can be split to have a more specific classification.
- `min_samples_leaf`
 - The `min_samples_leaf` parameter specifies the minimum number of samples required to be at a leaf node. The default value for this parameter is 1, which means that every leaf must have at least 1 sample that it classifies.

```
param_test1 = {'n_estimators':range(10,71,10), {'max_depth':range(3,14,2),
'min_samples_split':range(50,201,20), 'min_samples_leaf':range(30,60,10) } gsearch1 =
GridSearchCV(estimator = RandomForestClassifier(min_samples_split=100
min_samples_leaf=20,max_depth=8,max_features=' param_grid = param_test1,
scoring='roc_auc',cv=5) gsearch1.fit(X,y) gsearch1.bestparams, gsearch1.bestscore
```

XGBoost

XGBoost is a highly optimized framework for gradient boosting, an algorithm that iteratively combines the predictions of several weak learners such as decision trees to produce a much stronger and more robust model.

In [193...

```

from scipy import stats
from sklearn.metrics import (
    accuracy_score, confusion_matrix, classification_report,
    roc_auc_score, roc_curve, auc,
    plot_confusion_matrix, plot_roc_curve
)
from sklearn.model_selection import train_test_split, RandomizedSearchCV
from xgboost import XGBClassifier

```

```

param_grid = dict( n_estimators=stats.randint(10, 500), max_depth=stats.randint(1, 10),
learning_rate=stats.uniform(0, 1) )

```

```

xgb_clf = XGBClassifier() xgb_cv = RandomizedSearchCV( xgb_clf, param_grid, cv=3,
n_iter=60, scoring='roc_auc', n_jobs=-1, verbose=1) xgb_cv.fit(X_train, y_train)

best_params = xgb_cv.bestparams print(best_params)

```

In [175...

```

best_params['booster'] = 'gblinear'
print(f"Best Parameters: {best_params}")

```

```

Best Parameters: {'learning_rate': 0.0012592383320760847, 'max_depth': 8, 'n_e
stimators': 204, 'booster': 'gblinear'}

```

In [212...

```

xgb_clf = XGBClassifier({'learning_rate': 0.0012592383320760847, 'max_depth':

xgb_clf.fit(train_features, train_labels)

y_test_pred = xgb_clf.predict(test_features)

```

```

/Users/yanzhenlei/opt/anaconda3/lib/python3.8/site-packages/xgboost/core.py:41
6: FutureWarning: Pass `objective` as keyword args.  Passing these as position
al arguments will be considered as error in future releases.

```

```

warnings.warn(
/Users/yanzhenlei/opt/anaconda3/lib/python3.8/site-packages/xgboost/sklearn.py
:888: UserWarning: The use of label encoder in XGBClassifier is deprecated and
will be removed in a future release. To remove this warning, do the following:
1) Pass option use_label_encoder=False when constructing XGBClassifier object;
and 2) Encode your labels (y) as integers starting with 0, i.e. 0, 1, 2, ...,
[num_class - 1].

```

```

warnings.warn(label_encoder_deprecation_msg, UserWarning)
[19:54:39] WARNING: /opt/concourse/worker/volumes/live/7a2b9f41-3287-451b-6691
-43e9a6c0910f/volume/xgboost-split_1619728204606/work/src/learner.cc:1061: Sta
rting in XGBoost 1.3.0, the default evaluation metric used with the objective
'multi:softprob' was changed from 'merror' to 'mlogloss'. Explicitly set eval_
metric if you'd like to restore the old behavior.

```

In [213...

```
print("Mean squared error: %.2f" % mean_squared_error(test_labels, y_test_pre)
print('Mean Absolute Percentage Error (MAPE):', round(metrics.mean_absolute_p
print('Accuracy:', round(100*(1 - metrics.mean_absolute_percentage_error(test
```

Mean squared error: 0.62
Mean Absolute Percentage Error (MAPE): 4.64
Accuracy: 95.36

In [204...

```
# define model evaluation method
cv = RepeatedKfold(n_splits=5, n_repeats=3, random_state=1)
# evaluate model
scores = cross_val_score(xgb_clf, train_features, train_labels, scoring='neg_
print('Mean MAE: %.3f (%.3f)' % (scores.mean(), scores.std()) )
```

Mean MAE: -0.491 (0.042)

In [203...

```
# define model evaluation method
cv = RepeatedKfold(n_splits=5, n_repeats=3, random_state=1)
# evaluate model
scores = cross_val_score(rf, train_features, train_labels, scoring='neg_mean_
print('Mean MAE: %.3f (%.3f)' % (scores.mean(), scores.std()) )
```

Mean MAE: -0.383 (0.015)

Conclusion

- EDA
 - 10000 sample size with 55 columns.
 - Many variables containing outliers and missing values
 - many variables are high-imbalanced.
 - Interest rate distribution are right-skewed. If we use linear regression, we should log-transform the interest rate
 - Grades and subgrades are highly correlated to interest rate
- Model Selection
 - Randomforest Model has mean MAE-0.383 and MAPE 3.55
 - XGBoost Model has mean MAE -0.491 and MAPE 4.64
 - Randomforest would be a better choice
- Feature Selection
 - The feature I choose are basily about Grade and Subgrade -grade: Grade associated with the loan. -sub_grade: Detailed grade associated with the loan.
 - However, we don't know what does grade are given. Only when we find out what influence grades, we can deep dive into different variables that affecting interest rate.
- Next step:
 - Add more models(Neural Networks and Linear regression)
 - Explore more about how does grades and sub-grades influences the interest rate. Correlaiton does not mean causal inferences
 - Explore more on the parameters,optimizing the performance of the model

In []: